

assignment-aml-cats

April 1, 2024

Group - 03 Assignment by Gaurav Kudeshia & Anurodh Singh

Assignment-2: Exploring Convolutional Neural Networks (CNNs)

The primary aim of this assignment is to explore the capabilities of Convolutional Neural Networks in identifying and classifying objects within images. Students will engage in practical application of CNNs on image datasets, leveraging pre-trained models to evaluate their effectiveness and performance.

```
[1]: !unzip -/fs/ess/PGS0341/BA_64061_KSU_SEC1/data/dogs-vs-cats.zip
```

UnZip 6.00 of 20 April 2009, by Info-ZIP. Maintained by C. Spieler. Send bug reports using <http://www.info-zip.org/zip-bug.html>; see README for details.

Usage: unzip [-Z] [-opts[modifiers]] file[.zip] [list] [-x xlist] [-d exdir]
Default action is to extract files in list, except those in xlist, to exdir;
file[.zip] may be a wildcard. -Z => ZipInfo mode ("unzip -Z" for usage).

-p	extract files to pipe, no messages	-l	list files (short format)
-f	freshen existing files, create none	-t	test compressed archive data
-u	update files, create if necessary	-z	display archive comment only
-v	list verbosely/show version info	-T	timestamp archive to latest
-x	exclude files that follow (in xlist)	-d	extract files into exdir

modifiers:

-n	never overwrite existing files	-q	quiet mode (-qq => quieter)
-o	overwrite files WITHOUT prompting	-a	auto-convert any text files
-j	junk paths (do not make directories)	-aa	treat ALL files as text
-U	use escapes for all non-ASCII Unicode	-UU	ignore any Unicode fields
-C	match filenames case-insensitively	-L	make (some) names lowercase
-X	restore UID/GID info	-V	retain VMS version numbers
-K	keep setuid/setgid/tacky permissions	-M	pipe through "more" pager
-O CHARSET	specify a character encoding for DOS, Windows and OS/2 archives		
-I CHARSET	specify a character encoding for UNIX and other archives		

See "unzip -hh" or unzip.txt for more help. Examples:

```
unzip data1 -x joe    => extract all files except joe from zipfile data1.zip
unzip -p foo | more   => send contents of foo.zip via pipe into program more
unzip -fo foo ReadMe => quietly replace existing ReadMe if archive file newer
```

```
[2]: !unzip -o -qq /fs/ess/PGS0341/BA_64061_KSU_SEC1/data/dogs-vs-cats.zip
```

```
[3]: !unzip -qq/fs/ess/PGS0341/BA_64061_KSU_SEC1/data/dogs-vs-cats train.zip
```

UnZip 6.00 of 20 April 2009, by Info-ZIP. Maintained by C. Spieler. Send bug reports using <http://www.info-zip.org/zip-bug.html>; see README for details.

Usage: unzip [-Z] [-opts[modifiers]] file[.zip] [list] [-x xlist] [-d exdir]
Default action is to extract files in list, except those in xlist, to exdir;
file[.zip] may be a wildcard. -Z => ZipInfo mode ("unzip -Z" for usage).

-p	extract files to pipe, no messages	-l	list files (short format)
-f	freshen existing files, create none	-t	test compressed archive data
-u	update files, create if necessary	-z	display archive comment only
-v	list verbosely/show version info	-T	timestamp archive to latest
-x	exclude files that follow (in xlist)	-d	extract files into exdir

modifiers:

-n	never overwrite existing files	-q	quiet mode (-qq => quieter)
-o	overwrite files WITHOUT prompting	-a	auto-convert any text files
-j	junk paths (do not make directories)	-aa	treat ALL files as text
-U	use escapes for all non-ASCII Unicode	-UU	ignore any Unicode fields
-C	match filenames case-insensitively	-L	make (some) names lowercase
-X	restore UID/GID info	-V	retain VMS version numbers
-K	keep setuid/setgid/tacky permissions	-M	pipe through "more" pager
-O CHARSET	specify a character encoding for DOS, Windows and OS/2 archives		
-I CHARSET	specify a character encoding for UNIX and other archives		

See "unzip -hh" or unzip.txt for more help. Examples:

```
unzip data1 -x joe    => extract all files except joe from zipfile data1.zip
unzip -p foo | more   => send contents of foo.zip via pipe into program more
unzip -fo foo ReadMe => quietly replace existing ReadMe if archive file newer
```

Q1, Consider the Cats & Dogs example. Start initially with a training sample of 1000, a validation sample of 500, and a test sample of 500 (like in the text). Use any technique to reduce overfitting and improve performance in developing a network that you train from scratch. What performance did you achieve?

We're going to build a convolutional neural network from the ground up. Having loaded our dataset, it's time to split it into distinct subsets for training, validation, and testing. For this project, we'll allocate 1,000 images for training purposes, 500 images for validation, and another 500 images will be used for the test set.

Organizing the Dataset into Training, Validation, and Testing Groups

```
[4]: import os, shutil, pathlib

original_dir = pathlib.Path("train")
new_base_dir = pathlib.Path("cats_vs_dogs_small")
```

```
def make_subset(subset_name, start_index, end_index):
    for category in ("cat", "dog"):
        dir = new_base_dir / subset_name / category
        os.makedirs(dir, exist_ok=True)
        fnames = [f"{category}.{i}.jpg" for i in range(start_index, end_index)]
        for fname in fnames:
            src = original_dir / fname
            dst = dir / fname
            shutil.copyfile(src, dst)

make_subset("train", start_index=0, end_index=1000)
make_subset("validation", start_index=1000, end_index=1500)
make_subset("test", start_index=1500, end_index=2000)
```

0.0.1 Constructing the Model

For this network, we begin by processing images, represented as 3D tensors, which are initially reshaped. The process involves applying convolution operations using a 3x3 window (referred to as `kernel_size`), followed by max pooling operations with a 2x2 window (known as `pool_size`).

The objective of this task is to categorize the images into two classes: “cat” or “dog”. To achieve this, the architecture incorporates a dense layer towards the end, which plays a crucial role in determining the classification of the output as either “cat” or “dog”. This classification is facilitated by a single output node corresponding to these categories. Before reaching the dense layer, it’s necessary to transform the 3D tensor structure into a 1D format, a step accomplished by introducing a flattening layer.

Creating a Compact Convolutional Network for Cat vs. Dog Classification

```
[5]: from tensorflow import keras
    from tensorflow.keras import layers

    inputs = keras.Input(shape=(180, 180, 3))
    x = layers.Rescaling(1./255)(inputs)
    x = layers.Conv2D(filters=32, kernel_size=3, activation="relu")(x)
    x = layers.MaxPooling2D(pool_size=2)(x)
    x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)
    x = layers.MaxPooling2D(pool_size=2)(x)
    x = layers.Conv2D(filters=128, kernel_size=3, activation="relu")(x)
    x = layers.MaxPooling2D(pool_size=2)(x)
    x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
    x = layers.MaxPooling2D(pool_size=2)(x)
    x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
    x = layers.Flatten()(x)
    outputs = layers.Dense(1, activation="sigmoid")(x)
    model = keras.Model(inputs=inputs, outputs=outputs)
```

```
[6]: model.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 180, 180, 3)]	0
rescaling (Rescaling)	(None, 180, 180, 3)	0
conv2d (Conv2D)	(None, 178, 178, 32)	896
max_pooling2d (MaxPooling2D)	(None, 89, 89, 32)	0
conv2d_1 (Conv2D)	(None, 87, 87, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 43, 43, 64)	0
conv2d_2 (Conv2D)	(None, 41, 41, 128)	73856
max_pooling2d_2 (MaxPooling2D)	(None, 20, 20, 128)	0
conv2d_3 (Conv2D)	(None, 18, 18, 256)	295168
max_pooling2d_3 (MaxPooling2D)	(None, 9, 9, 256)	0
conv2d_4 (Conv2D)	(None, 7, 7, 256)	590080
flatten (Flatten)	(None, 12544)	0
dense (Dense)	(None, 1)	12545
Total params: 991,041		
Trainable params: 991,041		
Non-trainable params: 0		

Configuring the model for training

```
[7]: model.compile(loss="binary_crossentropy",  
                  optimizer="rmsprop",  
                  metrics=["accuracy"])
```

0.0.2 Data preprocessing

Using `image_dataset_from_directory` to read images

```
[8]: from tensorflow.keras.utils import image_dataset_from_directory  
  
train_dataset = image_dataset_from_directory(
```

```

    new_base_dir / "train",
    image_size=(180, 180),
    batch_size=32)
validation_dataset = image_dataset_from_directory(
    new_base_dir / "validation",
    image_size=(180, 180),
    batch_size=32)
test_dataset = image_dataset_from_directory(
    new_base_dir / "test",
    image_size=(180, 180),
    batch_size=32)

```

Found 2000 files belonging to 2 classes.
 Found 1000 files belonging to 2 classes.
 Found 2000 files belonging to 2 classes.

```

[9]: import numpy as np
import tensorflow as tf
random_numbers = np.random.normal(size=(1000, 16))
dataset = tf.data.Dataset.from_tensor_slices(random_numbers)

```

```

[10]: for i, element in enumerate(dataset):
    print(element.shape)
    if i >= 2:
        break

```

(16,)
 (16,)
 (16,)

```

[11]: batched_dataset = dataset.batch(32)
for i, element in enumerate(batched_dataset):
    print(element.shape)
    if i >= 2:
        break

```

(32, 16)
 (32, 16)
 (32, 16)

```

[12]: reshaped_dataset = dataset.map(lambda x: tf.reshape(x, (4, 4)))
for i, element in enumerate(reshaped_dataset):
    print(element.shape)
    if i >= 2:
        break

```

(4, 4)
 (4, 4)

(4, 4)

Showing the Dimensions of Data and Labels Provided by the Dataset

```
[13]: for data_batch, labels_batch in train_dataset:
      print("data batch shape:", data_batch.shape)
      print("labels batch shape:", labels_batch.shape)
      break
```

data batch shape: (32, 180, 180, 3)

labels batch shape: (32,)

Training the Model with a Dataset

```
[14]: callbacks = [
      keras.callbacks.ModelCheckpoint(
          filepath="convnet_from_scratch.keras",
          save_best_only=True,
          monitor="val_loss")
      ]
      history = model.fit(
          train_dataset,
          epochs=10,
          validation_data=validation_dataset,
          callbacks=callbacks)
```

Epoch 1/10

63/63 [=====] - 46s 726ms/step - loss: 0.7247 - accuracy: 0.5230 - val_loss: 0.6839 - val_accuracy: 0.5840

Epoch 2/10

63/63 [=====] - 45s 711ms/step - loss: 0.7062 - accuracy: 0.5555 - val_loss: 0.9294 - val_accuracy: 0.5070

Epoch 3/10

63/63 [=====] - 45s 713ms/step - loss: 0.6871 - accuracy: 0.6100 - val_loss: 0.6607 - val_accuracy: 0.6090

Epoch 4/10

63/63 [=====] - 45s 713ms/step - loss: 0.6377 - accuracy: 0.6505 - val_loss: 0.6445 - val_accuracy: 0.6060

Epoch 5/10

63/63 [=====] - 45s 713ms/step - loss: 0.5790 - accuracy: 0.6920 - val_loss: 0.8312 - val_accuracy: 0.6310

Epoch 6/10

63/63 [=====] - 45s 714ms/step - loss: 0.5475 - accuracy: 0.7215 - val_loss: 0.5828 - val_accuracy: 0.6970

Epoch 7/10

63/63 [=====] - 45s 708ms/step - loss: 0.5099 - accuracy: 0.7525 - val_loss: 0.6031 - val_accuracy: 0.6910

Epoch 8/10

63/63 [=====] - 45s 709ms/step - loss: 0.4653 -

accuracy: 0.7915 - val_loss: 0.8097 - val_accuracy: 0.6990

Epoch 9/10

63/63 [=====] - 45s 706ms/step - loss: 0.4182 -

accuracy: 0.8110 - val_loss: 0.5778 - val_accuracy: 0.7350


Epoch 10/10

63/63 [=====] - 45s 714ms/step - loss: 0.3663 -

accuracy: 0.8375 - val_loss: 0.5970 - val_accuracy: 0.7400

Visualizing Training Loss and Accuracy Curves

```
[15]: import matplotlib.pyplot as plt
accuracy = history.history["accuracy"]
val_accuracy = history.history["val_accuracy"]
loss = history.history["loss"]
val_loss = history.history["val_loss"]
epochs = range(1, len(accuracy) + 1)
plt.plot(epochs, accuracy, "bo", label="Training accuracy")
plt.plot(epochs, val_accuracy, "b", label="Validation accuracy")
plt.title("Training and validation accuracy")
plt.legend()
plt.figure()
plt.plot(epochs, loss, "bo", label="Training loss")
plt.plot(epochs, val_loss, "b", label="Validation loss")
plt.title("Training and validation loss")
plt.legend()
plt.show()
```





Assessing the model's performance using the test dataset.

```
[16]: test_model = keras.models.load_model("convnet_from_scratch.keras")
test_loss, test_acc = test_model.evaluate(test_dataset)
print(f"Test accuracy: {test_acc:.3f}")
```

```
63/63 [=====] - 10s 160ms/step - loss: 0.6221 -
accuracy: 0.7300
Test accuracy: 0.730
```

The model's validation and test accuracy currently stands at a relatively low **73%**. - To enhance the model's performance, we plan to implement several strategies, including:

- a) Data Augmentation
- b) Dropout Technique
- c) A combination of Data Augmentation and the Dropout Technique

0.0.3 a, By employing data augmentation

Implement a Data Augmentation Stage for an Image Processing Model










```
[17]: data_augmentation = keras.Sequential(
    [
        layers.RandomFlip("horizontal"),
```



```
layers.RandomRotation(0.1),  
layers.RandomZoom(0.2),  
]  
)
```

Displaying some randomly augmented training images

```
[18]: plt.figure(figsize=(10, 10))  
for images, _ in train_dataset.take(1):  
    for i in range(9):  
        augmented_images = data_augmentation(images)  
        ax = plt.subplot(3, 3, i + 1)  
        plt.imshow(augmented_images[0].numpy().astype("uint8"))  
        plt.axis("off")
```



Q2. Increasing Training sample size to 2870 samples

```
[19]: import os, shutil, pathlib

original_dir = pathlib.Path("train")
new_base_dir = pathlib.Path("cats_vs_dogs_small_01")

def make_subset(subset_name, start_index, end_index):
    for category in ("cat", "dog"):
        dir = new_base_dir / subset_name / category
        os.makedirs(dir, exist_ok=True)
        fnames = [f"{category}.{i}.jpg" for i in range(start_index, end_index)]
        for fname in fnames:
            src = original_dir / fname
            dst = dir / fname
            shutil.copyfile(src, dst)

make_subset("train", start_index=0, end_index=1870)
make_subset("validation", start_index=1870, end_index=2370)
make_subset("test", start_index=2370, end_index=2870)
```

Defining a new convnet that includes image augmentation and dropout

```
[20]: inputs = keras.Input(shape=(180, 180, 3))
x = data_augmentation(inputs)
x = layers.Rescaling(1./255)(x)
x = layers.Conv2D(filters=32, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=128, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.Flatten()(x)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs=inputs, outputs=outputs)

model.compile(loss="binary_crossentropy",
              optimizer="rmsprop",
              metrics=["accuracy"])
```

Training the regularized convnet

```
[22]: callbacks = [
    keras.callbacks.ModelCheckpoint(
```

```

        filepath="convnet_from_scratch_with_augmentation.keras",
        save_best_only=True,
        monitor="val_loss")
]
history = model.fit(
    train_dataset,
    epochs=45,
    validation_data=validation_dataset,
    callbacks=callbacks)

```

Epoch 1/45

63/63 [=====] - 47s 752ms/step - loss: 0.6935 - accuracy: 0.5275 - val_loss: 0.6748 - val_accuracy: 0.5590

Epoch 2/45

63/63 [=====] - 47s 742ms/step - loss: 0.6701 - accuracy: 0.6140 - val_loss: 0.6567 - val_accuracy: 0.6130

Epoch 3/45

63/63 [=====] - 47s 745ms/step - loss: 0.6581 - accuracy: 0.6535 - val_loss: 0.6519 - val_accuracy: 0.5830

Epoch 4/45

63/63 [=====] - 47s 747ms/step - loss: 0.6348 - accuracy: 0.6590 - val_loss: 0.6216 - val_accuracy: 0.6320

Epoch 5/45

63/63 [=====] - 47s 747ms/step - loss: 0.6437 - accuracy: 0.6595 - val_loss: 0.6396 - val_accuracy: 0.6200

Epoch 6/45

63/63 [=====] - 47s 747ms/step - loss: 0.6065 - accuracy: 0.6735 - val_loss: 0.6303 - val_accuracy: 0.6260

Epoch 7/45

63/63 [=====] - 47s 747ms/step - loss: 0.5992 - accuracy: 0.6690 - val_loss: 0.5956 - val_accuracy: 0.6630

Epoch 8/45

63/63 [=====] - 47s 749ms/step - loss: 0.5961 - accuracy: 0.6875 - val_loss: 0.5844 - val_accuracy: 0.6910

Epoch 9/45

63/63 [=====] - 47s 748ms/step - loss: 0.5702 - accuracy: 0.7085 - val_loss: 0.5684 - val_accuracy: 0.6940

Epoch 10/45

63/63 [=====] - 47s 748ms/step - loss: 0.5564 - accuracy: 0.7190 - val_loss: 0.6759 - val_accuracy: 0.6940

Epoch 11/45

63/63 [=====] - 47s 748ms/step - loss: 0.5551 - accuracy: 0.7145 - val_loss: 0.5375 - val_accuracy: 0.7350

Epoch 12/45

63/63 [=====] - 47s 747ms/step - loss: 0.5614 - accuracy: 0.7240 - val_loss: 0.5541 - val_accuracy: 0.7200

Epoch 13/45

63/63 [=====] - 47s 748ms/step - loss: 0.5202 -
accuracy: 0.7625 - val_loss: 0.5239 - val_accuracy: 0.7460
Epoch 14/45
63/63 [=====] - 47s 740ms/step - loss: 0.5167 -
accuracy: 0.7505 - val_loss: 0.6480 - val_accuracy: 0.6910
Epoch 15/45
63/63 [=====] - 47s 747ms/step - loss: 0.5192 -
accuracy: 0.7450 - val_loss: 0.5331 - val_accuracy: 0.7320
Epoch 16/45
63/63 [=====] - 47s 743ms/step - loss: 0.4957 -
accuracy: 0.7645 - val_loss: 0.5934 - val_accuracy: 0.7250
Epoch 17/45
63/63 [=====] - 47s 745ms/step - loss: 0.4875 -
accuracy: 0.7605 - val_loss: 0.6310 - val_accuracy: 0.7470
Epoch 18/45
63/63 [=====] - 47s 750ms/step - loss: 0.4869 -
accuracy: 0.7750 - val_loss: 0.5260 - val_accuracy: 0.7760
Epoch 19/45
63/63 [=====] - 47s 745ms/step - loss: 0.4682 -
accuracy: 0.7845 - val_loss: 0.4882 - val_accuracy: 0.7840
Epoch 20/45
63/63 [=====] - 47s 743ms/step - loss: 0.4687 -
accuracy: 0.7785 - val_loss: 0.4578 - val_accuracy: 0.7940
Epoch 21/45
63/63 [=====] - 47s 747ms/step - loss: 0.4672 -
accuracy: 0.7850 - val_loss: 0.4628 - val_accuracy: 0.7920
Epoch 22/45
63/63 [=====] - 47s 747ms/step - loss: 0.4405 -
accuracy: 0.8045 - val_loss: 0.4707 - val_accuracy: 0.7860
Epoch 23/45
63/63 [=====] - 47s 750ms/step - loss: 0.4471 -
accuracy: 0.7800 - val_loss: 0.4806 - val_accuracy: 0.7860
Epoch 24/45
63/63 [=====] - 47s 742ms/step - loss: 0.4304 -
accuracy: 0.8065 - val_loss: 0.7241 - val_accuracy: 0.6730
Epoch 25/45
63/63 [=====] - 47s 747ms/step - loss: 0.4344 -
accuracy: 0.7920 - val_loss: 0.4272 - val_accuracy: 0.8180
Epoch 26/45
63/63 [=====] - 47s 746ms/step - loss: 0.4119 -
accuracy: 0.8170 - val_loss: 0.4968 - val_accuracy: 0.7980
Epoch 27/45
63/63 [=====] - 47s 748ms/step - loss: 0.4060 -
accuracy: 0.8140 - val_loss: 0.4344 - val_accuracy: 0.8150
Epoch 28/45
63/63 [=====] - 47s 746ms/step - loss: 0.4196 -
accuracy: 0.8070 - val_loss: 0.4894 - val_accuracy: 0.7980
Epoch 29/45

63/63 [=====] - 47s 743ms/step - loss: 0.4063 -
accuracy: 0.8240 - val_loss: 0.5169 - val_accuracy: 0.7650
Epoch 30/45
63/63 [=====] - 47s 748ms/step - loss: 0.3670 -
accuracy: 0.8360 - val_loss: 0.4758 - val_accuracy: 0.8190
Epoch 31/45
63/63 [=====] - 47s 747ms/step - loss: 0.3940 -
accuracy: 0.8235 - val_loss: 0.5819 - val_accuracy: 0.7850
Epoch 32/45
63/63 [=====] - 47s 743ms/step - loss: 0.3843 -
accuracy: 0.8310 - val_loss: 0.4293 - val_accuracy: 0.8140
Epoch 33/45
63/63 [=====] - 47s 746ms/step - loss: 0.3713 -
accuracy: 0.8325 - val_loss: 0.4447 - val_accuracy: 0.8090
Epoch 34/45
63/63 [=====] - 47s 745ms/step - loss: 0.3788 -
accuracy: 0.8385 - val_loss: 0.5805 - val_accuracy: 0.7680
Epoch 35/45
63/63 [=====] - 47s 750ms/step - loss: 0.3473 -
accuracy: 0.8465 - val_loss: 0.4287 - val_accuracy: 0.8280
Epoch 36/45
63/63 [=====] - 47s 748ms/step - loss: 0.3319 -
accuracy: 0.8480 - val_loss: 0.4798 - val_accuracy: 0.8050
Epoch 37/45
63/63 [=====] - 47s 747ms/step - loss: 0.3308 -
accuracy: 0.8600 - val_loss: 0.3999 - val_accuracy: 0.8290
Epoch 38/45
63/63 [=====] - 47s 745ms/step - loss: 0.3331 -
accuracy: 0.8590 - val_loss: 0.5823 - val_accuracy: 0.7700
Epoch 39/45
63/63 [=====] - 47s 750ms/step - loss: 0.3450 -
accuracy: 0.8580 - val_loss: 0.4573 - val_accuracy: 0.8190
Epoch 40/45
63/63 [=====] - 47s 745ms/step - loss: 0.3291 -
accuracy: 0.8655 - val_loss: 0.4326 - val_accuracy: 0.8270
Epoch 41/45
63/63 [=====] - 47s 748ms/step - loss: 0.3089 -
accuracy: 0.8770 - val_loss: 0.4746 - val_accuracy: 0.8000
Epoch 42/45
63/63 [=====] - 47s 745ms/step - loss: 0.2976 -
accuracy: 0.8730 - val_loss: 0.5727 - val_accuracy: 0.8080
Epoch 43/45
63/63 [=====] - 47s 746ms/step - loss: 0.3086 -
accuracy: 0.8775 - val_loss: 0.4007 - val_accuracy: 0.8400
Epoch 44/45
63/63 [=====] - 47s 742ms/step - loss: 0.3114 -
accuracy: 0.8755 - val_loss: 0.4103 - val_accuracy: 0.8260
Epoch 45/45

```
63/63 [=====] - 47s 745ms/step - loss: 0.2869 -  
accuracy: 0.8855 - val_loss: 0.4800 - val_accuracy: 0.8250
```

Assessment of the Model Using the Test Data

```
[23]: test_model = keras.models.load_model(  
        "convnet_from_scratch_with_augmentation.keras")  
test_loss, test_acc = test_model.evaluate(test_dataset)  
print(f"Test accuracy: {test_acc:.3f}")
```

```
63/63 [=====] - 10s 159ms/step - loss: 0.4255 -  
accuracy: 0.8135  
Test accuracy: 0.813
```

Q3. Increasing Training sample size to 3100 samples

The number of training samples has been expanded to 3100, and the impact on model performance is discussed in the provided summary.

```
[24]: import os, shutil, pathlib  
  
original_dir = pathlib.Path("train")  
new_base_dir = pathlib.Path("cats_vs_dogs_small_02")  
  
def make_subset(subset_name, start_index, end_index):  
    for category in ("cat", "dog"):  
        dir = new_base_dir / subset_name / category  
        os.makedirs(dir, exist_ok=True)  
        fnames = [f"{category}.{i}.jpg" for i in range(start_index, end_index)]  
        for fname in fnames:  
            src = original_dir / fname  
            dst = dir / fname  
            shutil.copyfile(src, dst)  
  
make_subset("train", start_index=0, end_index=2100)  
make_subset("validation", start_index=2100, end_index=2600)  
make_subset("test", start_index=2600, end_index=3100)
```

Defining a new convnet that includes image augmentation and dropout

```
[25]: inputs = keras.Input(shape=(180, 180, 3))  
x = data_augmentation(inputs)  
x = layers.Rescaling(1./255)(x)  
x = layers.Conv2D(filters=32, kernel_size=3, activation="relu")(x)  
x = layers.MaxPooling2D(pool_size=2)(x)  
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)  
x = layers.MaxPooling2D(pool_size=2)(x)  
x = layers.Conv2D(filters=128, kernel_size=3, activation="relu")(x)  
x = layers.MaxPooling2D(pool_size=2)(x)  
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
```

```

x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.Flatten()(x)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs=inputs, outputs=outputs)

model.compile(loss="binary_crossentropy",
              optimizer="rmsprop",
              metrics=["accuracy"])

```

Now, we proceed to train the convnet that has been regularized.

```

[26]: callbacks = [
        keras.callbacks.ModelCheckpoint(
            filepath="convnet_from_scratch_with_augmentation.keras",
            save_best_only=True,
            monitor="val_loss")
    ]
history = model.fit(
    train_dataset,
    epochs=20,
    validation_data=validation_dataset,
    callbacks=callbacks)

```

Epoch 1/20

```
63/63 [=====] - 49s 755ms/step - loss: 0.7548 -
accuracy: 0.4920 - val_loss: 0.6919 - val_accuracy: 0.5030
```

Epoch 2/20

```
63/63 [=====] - 47s 751ms/step - loss: 0.6993 -
accuracy: 0.5115 - val_loss: 0.6876 - val_accuracy: 0.5360
```

Epoch 3/20

```
63/63 [=====] - 47s 748ms/step - loss: 0.6978 -
accuracy: 0.5535 - val_loss: 0.6811 - val_accuracy: 0.5560
```

Epoch 4/20

```
63/63 [=====] - 47s 750ms/step - loss: 0.6664 -
accuracy: 0.6055 - val_loss: 0.6350 - val_accuracy: 0.6570
```

Epoch 5/20

```
63/63 [=====] - 47s 742ms/step - loss: 0.6583 -
accuracy: 0.6355 - val_loss: 0.6771 - val_accuracy: 0.6010
```

Epoch 6/20

```
63/63 [=====] - 47s 744ms/step - loss: 0.6433 -
accuracy: 0.6460 - val_loss: 0.6111 - val_accuracy: 0.6590
```

Epoch 7/20

```
63/63 [=====] - 47s 744ms/step - loss: 0.6218 -
accuracy: 0.6550 - val_loss: 0.5932 - val_accuracy: 0.6990
```

Epoch 8/20

```
63/63 [=====] - 47s 748ms/step - loss: 0.5977 -
```

```

accuracy: 0.6840 - val_loss: 0.5715 - val_accuracy: 0.6990
Epoch 9/20
63/63 [=====] - 47s 745ms/step - loss: 0.5801 -
accuracy: 0.7030 - val_loss: 0.5786 - val_accuracy: 0.6920
Epoch 10/20
63/63 [=====] - 47s 744ms/step - loss: 0.5806 -
accuracy: 0.6990 - val_loss: 0.7779 - val_accuracy: 0.6100
Epoch 11/20
63/63 [=====] - 47s 745ms/step - loss: 0.5682 -
accuracy: 0.7140 - val_loss: 0.5766 - val_accuracy: 0.7150
Epoch 12/20
63/63 [=====] - 47s 743ms/step - loss: 0.5529 -
accuracy: 0.7375 - val_loss: 0.7651 - val_accuracy: 0.6420
Epoch 13/20
63/63 [=====] - 47s 747ms/step - loss: 0.5450 -
accuracy: 0.7335 - val_loss: 0.6488 - val_accuracy: 0.6840
Epoch 14/20
63/63 [=====] - 47s 750ms/step - loss: 0.5450 -
accuracy: 0.7350 - val_loss: 0.5149 - val_accuracy: 0.7420
Epoch 15/20
63/63 [=====] - 47s 747ms/step - loss: 0.5254 -
accuracy: 0.7395 - val_loss: 0.4938 - val_accuracy: 0.7630
Epoch 16/20
63/63 [=====] - 47s 744ms/step - loss: 0.5226 -
accuracy: 0.7390 - val_loss: 0.5729 - val_accuracy: 0.6760
Epoch 17/20
63/63 [=====] - 47s 752ms/step - loss: 0.5022 -
accuracy: 0.7630 - val_loss: 0.4732 - val_accuracy: 0.7790
Epoch 18/20
63/63 [=====] - 47s 739ms/step - loss: 0.4907 -
accuracy: 0.7725 - val_loss: 0.5743 - val_accuracy: 0.7230
Epoch 19/20
63/63 [=====] - 47s 746ms/step - loss: 0.4887 -
accuracy: 0.7730 - val_loss: 0.5555 - val_accuracy: 0.7540
Epoch 20/20
63/63 [=====] - 47s 747ms/step - loss: 0.4712 -
accuracy: 0.7770 - val_loss: 0.4672 - val_accuracy: 0.7980

```

Evaluating the test dataset

```

[27]: test_model = keras.models.load_model(
        "convnet_from_scratch_with_augmentation.keras")
test_loss, test_acc = test_model.evaluate(test_dataset)
print(f"Test accuracy: {test_acc:.3f}")

```

```

63/63 [=====] - 10s 160ms/step - loss: 0.5118 -
accuracy: 0.7660
Test accuracy: 0.766

```


Q4.Fine Tuning of the pretrained models

We will adjust the pretrained model by experimenting with various sizes of training samples and then assess its effectiveness based on the performance of the models we previously constructed.

Pre-Trained Model with 1000 Training Samples

Initializing the VGG16 convolutional base and setting it to a non-trainable state

Instantiating the VGG16 convolutional base

```
[28]: conv_base = keras.applications.vgg16.VGG16(  
      weights="imagenet",  
      include_top=False,  
      input_shape=(180, 180, 3))
```

```
[29]: conv_base.summary()
```

Model: "vgg16"

Layer (type)	Output Shape	Param #
input_4 (InputLayer)	[(None, 180, 180, 3)]	0
block1_conv1 (Conv2D)	(None, 180, 180, 64)	1792
block1_conv2 (Conv2D)	(None, 180, 180, 64)	36928
block1_pool (MaxPooling2D)	(None, 90, 90, 64)	0
block2_conv1 (Conv2D)	(None, 90, 90, 128)	73856
block2_conv2 (Conv2D)	(None, 90, 90, 128)	147584
block2_pool (MaxPooling2D)	(None, 45, 45, 128)	0
block3_conv1 (Conv2D)	(None, 45, 45, 256)	295168
block3_conv2 (Conv2D)	(None, 45, 45, 256)	590080
block3_conv3 (Conv2D)	(None, 45, 45, 256)	590080
block3_pool (MaxPooling2D)	(None, 22, 22, 256)	0
block4_conv1 (Conv2D)	(None, 22, 22, 512)	1180160
block4_conv2 (Conv2D)	(None, 22, 22, 512)	2359808
block4_conv3 (Conv2D)	(None, 22, 22, 512)	2359808

block4_pool (MaxPooling2D)	(None, 11, 11, 512)	0

block5_conv1 (Conv2D)	(None, 11, 11, 512)	2359808

block5_conv2 (Conv2D)	(None, 11, 11, 512)	2359808

block5_conv3 (Conv2D)	(None, 11, 11, 512)	2359808

block5_pool (MaxPooling2D)	(None, 5, 5, 512)	0
=====		
Total params: 14,714,688		
Trainable params: 14,714,688		
Non-trainable params: 0		

Fast feature extraction without data augmentation Extracting the VGG16 features and corresponding labels

```
[30]: import numpy as np

def get_features_and_labels(dataset):
    all_features = []
    all_labels = []
    for images, labels in dataset:
        preprocessed_images = keras.applications.vgg16.preprocess_input(images)
        features = conv_base.predict(preprocessed_images)
        all_features.append(features)
        all_labels.append(labels)
    return np.concatenate(all_features), np.concatenate(all_labels)

train_features, train_labels = get_features_and_labels(train_dataset)
val_features, val_labels = get_features_and_labels(validation_dataset)
test_features, test_labels = get_features_and_labels(test_dataset)
```

```
[31]: train_features.shape
```

```
[31]: (2000, 5, 5, 512)
```

Defining and training the densely connected classifier

```
[33]: inputs = keras.Input(shape=(5, 5, 512))
x = layers.Flatten()(inputs)
x = layers.Dense(256)(x)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs, outputs)
model.compile(loss="binary_crossentropy",
              optimizer="rmsprop",
```

```

        metrics=["accuracy"])

callbacks = [
    keras.callbacks.ModelCheckpoint(
        filepath="feature_extraction.keras",
        save_best_only=True,
        monitor="val_loss")
]
history = model.fit(
    train_features, train_labels,
    epochs=30,
    validation_data=(val_features, val_labels),
    callbacks=callbacks)

```

Epoch 1/30

63/63 [=====] - 2s 33ms/step - loss: 15.5443 -
accuracy: 0.9190 - val_loss: 5.5484 - val_accuracy: 0.9650

Epoch 2/30

63/63 [=====] - 2s 33ms/step - loss: 2.9176 - accuracy:
0.9765 - val_loss: 8.3800 - val_accuracy: 0.9630

Epoch 3/30

63/63 [=====] - 2s 33ms/step - loss: 1.9215 - accuracy:
0.9870 - val_loss: 4.2566 - val_accuracy: 0.9730

Epoch 4/30

63/63 [=====] - 2s 32ms/step - loss: 1.3324 - accuracy:
0.9885 - val_loss: 5.3442 - val_accuracy: 0.9680

Epoch 5/30

63/63 [=====] - 2s 33ms/step - loss: 0.6369 - accuracy:
0.9945 - val_loss: 5.4410 - val_accuracy: 0.9730

Epoch 6/30

63/63 [=====] - 2s 33ms/step - loss: 0.5982 - accuracy:
0.9950 - val_loss: 4.4904 - val_accuracy: 0.9730

Epoch 7/30

63/63 [=====] - 2s 33ms/step - loss: 0.7402 - accuracy:
0.9935 - val_loss: 3.9737 - val_accuracy: 0.9790

Epoch 8/30

63/63 [=====] - 2s 32ms/step - loss: 0.1541 - accuracy:
0.9960 - val_loss: 3.8200 - val_accuracy: 0.9760

Epoch 9/30

63/63 [=====] - 2s 32ms/step - loss: 0.2506 - accuracy:
0.9965 - val_loss: 3.9778 - val_accuracy: 0.9770

Epoch 10/30

63/63 [=====] - 2s 33ms/step - loss: 0.0647 - accuracy:
0.9990 - val_loss: 3.3827 - val_accuracy: 0.9770

Epoch 11/30

63/63 [=====] - 2s 33ms/step - loss: 0.4268 - accuracy:
0.9965 - val_loss: 6.7100 - val_accuracy: 0.9660

Epoch 12/30
63/63 [=====] - 2s 33ms/step - loss: 0.2706 - accuracy: 0.9970 - val_loss: 4.8512 - val_accuracy: 0.9730

Epoch 13/30
63/63 [=====] - 2s 33ms/step - loss: 0.1408 - accuracy: 0.9975 - val_loss: 4.6766 - val_accuracy: 0.9780

Epoch 14/30
63/63 [=====] - 2s 32ms/step - loss: 0.1350 - accuracy: 0.9990 - val_loss: 4.6217 - val_accuracy: 0.9820

Epoch 15/30
63/63 [=====] - 2s 33ms/step - loss: 0.3154 - accuracy: 0.9980 - val_loss: 4.4168 - val_accuracy: 0.9760

Epoch 16/30
63/63 [=====] - 2s 33ms/step - loss: 2.0408e-31 - accuracy: 1.0000 - val_loss: 4.4168 - val_accuracy: 0.9760

Epoch 17/30
63/63 [=====] - 2s 33ms/step - loss: 1.4306e-15 - accuracy: 1.0000 - val_loss: 4.4168 - val_accuracy: 0.9760

Epoch 18/30
63/63 [=====] - 2s 33ms/step - loss: 0.0000e+00 - accuracy: 1.0000 - val_loss: 4.4168 - val_accuracy: 0.9760

Epoch 19/30
63/63 [=====] - 2s 33ms/step - loss: 0.0043 - accuracy: 0.9990 - val_loss: 4.1333 - val_accuracy: 0.9810

Epoch 20/30
63/63 [=====] - 2s 33ms/step - loss: 0.1289 - accuracy: 0.9985 - val_loss: 4.3503 - val_accuracy: 0.9810

Epoch 21/30
63/63 [=====] - 2s 33ms/step - loss: 0.0066 - accuracy: 0.9990 - val_loss: 4.4840 - val_accuracy: 0.9790

Epoch 22/30
63/63 [=====] - 2s 33ms/step - loss: 0.0709 - accuracy: 0.9990 - val_loss: 4.3030 - val_accuracy: 0.9790

Epoch 23/30
63/63 [=====] - 2s 33ms/step - loss: 0.1636 - accuracy: 0.9980 - val_loss: 5.5896 - val_accuracy: 0.9720

Epoch 24/30
63/63 [=====] - 2s 33ms/step - loss: 0.0022 - accuracy: 0.9995 - val_loss: 5.9060 - val_accuracy: 0.9690

Epoch 25/30
63/63 [=====] - 2s 33ms/step - loss: 0.0603 - accuracy: 0.9990 - val_loss: 5.0379 - val_accuracy: 0.9760

Epoch 26/30
63/63 [=====] - 2s 33ms/step - loss: 4.9801e-09 - accuracy: 1.0000 - val_loss: 5.5193 - val_accuracy: 0.9750

Epoch 27/30
63/63 [=====] - 2s 33ms/step - loss: 0.0435 - accuracy: 0.9995 - val_loss: 8.3777 - val_accuracy: 0.9680

```

Epoch 28/30
63/63 [=====] - 2s 33ms/step - loss: 0.0068 - accuracy:
0.9995 - val_loss: 5.2382 - val_accuracy: 0.9780
Epoch 29/30
63/63 [=====] - 2s 33ms/step - loss: 0.2427 - accuracy:
0.9990 - val_loss: 5.4280 - val_accuracy: 0.9770
Epoch 30/30
63/63 [=====] - 2s 33ms/step - loss: 0.0633 - accuracy:
0.9990 - val_loss: 5.5203 - val_accuracy: 0.9790


```


Plotting the results

```

[34]: import matplotlib.pyplot as plt
acc = history.history["accuracy"]
val_acc = history.history["val_accuracy"]
loss = history.history["loss"]
val_loss = history.history["val_loss"]
epochs = range(1, len(acc) + 1)
plt.plot(epochs, acc, "bo", label="Training accuracy")
plt.plot(epochs, val_acc, "b", label="Validation accuracy")
plt.title("Training and validation accuracy")
plt.legend()
plt.figure()
plt.plot(epochs, loss, "bo", label="Training loss")
plt.plot(epochs, val_loss, "b", label="Validation loss")
plt.title("Training and validation loss")
plt.legend()
plt.show()

```





Feature extraction together with data augmentation Instantiating and freezing the VGG16 convolutional base

```
[35]: conv_base = keras.applications.vgg16.VGG16(
        weights="imagenet",
        include_top=False)
conv_base.trainable = False
```

Printing the list of trainable weights before and after freezing

```
[36]: conv_base.trainable = True
print("This is the number of trainable weights "
      "before freezing the conv base:", len(conv_base.trainable_weights))
```

This is the number of trainable weights before freezing the conv base: 26

```
[37]: conv_base.trainable = False
print("This is the number of trainable weights "
      "after freezing the conv base:", len(conv_base.trainable_weights))
```

This is the number of trainable weights after freezing the conv base: 0

Adding a data augmentation stage and a classifier to the convolutional base

```
[38]: data_augmentation = keras.Sequential(
    [
        layers.RandomFlip("horizontal"),
        layers.RandomRotation(0.1),
        layers.RandomZoom(0.2),
    ]
)

inputs = keras.Input(shape=(180, 180, 3))
x = data_augmentation(inputs)
x = keras.applications.vgg16.preprocess_input(x)
x = conv_base(x)
x = layers.Flatten()(x)
x = layers.Dense(256)(x)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs, outputs)
model.compile(loss="binary_crossentropy",
              optimizer="rmsprop",
              metrics=["accuracy"])
```

```
[39]: callbacks = [
    keras.callbacks.ModelCheckpoint(
        filepath="feature_extraction_with_data_augmentation.keras",
        save_best_only=True,
        monitor="val_loss")
]

history = model.fit(
    train_dataset,
    epochs=10,
    validation_data=validation_dataset,
    callbacks=callbacks)
```

```
Epoch 1/10
63/63 [=====] - 168s 3s/step - loss: 18.4987 -
accuracy: 0.8920 - val_loss: 14.4387 - val_accuracy: 0.9230
Epoch 2/10
63/63 [=====] - 167s 3s/step - loss: 6.1704 - accuracy:
0.9530 - val_loss: 8.7896 - val_accuracy: 0.9510
Epoch 3/10
63/63 [=====] - 167s 3s/step - loss: 6.2106 - accuracy:
0.9510 - val_loss: 4.6838 - val_accuracy: 0.9730
Epoch 4/10
63/63 [=====] - 167s 3s/step - loss: 5.0829 - accuracy:
0.9590 - val_loss: 7.3436 - val_accuracy: 0.9640
Epoch 5/10
63/63 [=====] - 167s 3s/step - loss: 4.5701 - accuracy:
```

```

0.9620 - val_loss: 4.5137 - val_accuracy: 0.9720
Epoch 6/10
63/63 [=====] - 167s 3s/step - loss: 3.9205 - accuracy:
0.9685 - val_loss: 3.1031 - val_accuracy: 0.9790
Epoch 7/10
63/63 [=====] - 167s 3s/step - loss: 3.7451 - accuracy:
0.9685 - val_loss: 4.0067 - val_accuracy: 0.9780
Epoch 8/10
63/63 [=====] - 166s 3s/step - loss: 3.0027 - accuracy:
0.9730 - val_loss: 6.3938 - val_accuracy: 0.9700
Epoch 9/10
63/63 [=====] - 167s 3s/step - loss: 3.2352 - accuracy:
0.9770 - val_loss: 4.9734 - val_accuracy: 0.9720
Epoch 10/10
63/63 [=====] - 166s 3s/step - loss: 2.5900 - accuracy:
0.9770 - val_loss: 4.2494 - val_accuracy: 0.9760

```

Evaluating the model on the test set

```

[ ]: test_model = keras.models.load_model(
    "feature_extraction_with_data_augmentation.keras")
test_loss, test_acc = test_model.evaluate(test_dataset)
print(f"Test accuracy: {test_acc:.3f}")

```

```

5/63 [=>...] - ETA: 1:41 - loss: 1.6594 - accuracy:
0.9875

```

0.0.4 Fine-tuning a pretrained model

```

[39]: conv_base.summary()

```

Model: "vgg16"

Layer (type)	Output Shape	Param #
input_6 (InputLayer)	[(None, None, None, 3)]	0
block1_conv1 (Conv2D)	(None, None, None, 64)	1792
block1_conv2 (Conv2D)	(None, None, None, 64)	36928
block1_pool (MaxPooling2D)	(None, None, None, 64)	0
block2_conv1 (Conv2D)	(None, None, None, 128)	73856
block2_conv2 (Conv2D)	(None, None, None, 128)	147584
block2_pool (MaxPooling2D)	(None, None, None, 128)	0


```

-----
block3_conv1 (Conv2D)          (None, None, None, 256)    295168
-----
block3_conv2 (Conv2D)          (None, None, None, 256)    590080
-----
block3_conv3 (Conv2D)          (None, None, None, 256)    590080
-----
block3_pool (MaxPooling2D)     (None, None, None, 256)    0
-----
block4_conv1 (Conv2D)          (None, None, None, 512)    1180160
-----
block4_conv2 (Conv2D)          (None, None, None, 512)    2359808
-----
block4_conv3 (Conv2D)          (None, None, None, 512)    2359808
-----
block4_pool (MaxPooling2D)     (None, None, None, 512)    0
-----
block5_conv1 (Conv2D)          (None, None, None, 512)    2359808
-----
block5_conv2 (Conv2D)          (None, None, None, 512)    2359808
-----
block5_conv3 (Conv2D)          (None, None, None, 512)    2359808
-----
block5_pool (MaxPooling2D)     (None, None, None, 512)    0
=====
Total params: 14,714,688
Trainable params: 0
Non-trainable params: 14,714,688
-----

```

Freezing all layers until the fourth from the last

```

[40]: conv_base.trainable = True
      for layer in conv_base.layers[:-4]:
          layer.trainable = False

```

Fine-tuning the model

```

[41]: model.compile(loss="binary_crossentropy",
                    optimizer=keras.optimizers.RMSprop(learning_rate=1e-5),
                    metrics=["accuracy"])

callbacks = [
    keras.callbacks.ModelCheckpoint(
        filepath="fine_tuning.keras",
        save_best_only=True,
        monitor="val_loss")
]

```

```

history = model.fit(
    train_dataset,
    epochs=10,
    validation_data=validation_dataset,
    callbacks=callbacks)

```

63/63 [=====] - 192s 3s/step - loss: 2.5842 - accuracy: 0.9730 - val_loss: 2.7219 - val_accuracy: 0.9780

```

[42]: model = keras.models.load_model("fine_tuning.keras")
test_loss, test_acc = model.evaluate(test_dataset)
print(f"Test accuracy: {test_acc:.3f}")

```

63/63 [=====] - 107s 2s/step - loss: 2.9272 - accuracy: 0.9735

Test accuracy: 0.974

```

[43]: conv_base = keras.applications.vgg16.VGG16(
    weights="imagenet",
    include_top=False,
    input_shape=(180, 180, 3))

```

```

[44]: conv_base = keras.applications.vgg16.VGG16(
    weights="imagenet",
    include_top=False)

conv_base.trainable = True
for layer in conv_base.layers[:-4]:
    layer.trainable = False

```

```

[46]: data_augmentation = keras.Sequential(
    [
        layers.RandomFlip("horizontal"),
        layers.RandomRotation(0.1),
        layers.RandomZoom(0.2),
    ]
)

inputs = keras.Input(shape=(180, 180, 3))
x = data_augmentation(inputs)
x = keras.applications.vgg16.preprocess_input(x)
x = conv_base(x)
x = layers.Flatten()(x)
x = layers.Dense(256)(x)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs, outputs)

```

```

model.compile(loss="binary_crossentropy",
              optimizer=keras.optimizers.RMSprop(learning_rate=1e-5),
              metrics=["accuracy"])

callbacks = [
    keras.callbacks.ModelCheckpoint(
        filepath="fine_tuning2.keras",
        save_best_only=True,
        monitor="val_loss")
]
history = model.fit(
    train_dataset,
    epochs=5,
    validation_data=validation_dataset,
    callbacks=callbacks)

```

63/63 [=====] - 193s 3s/step - loss: 3.6172 - accuracy: 0.7560 - val_loss: 0.5617 - val_accuracy: 0.9330

```

[47]: model = keras.models.load_model("fine_tuning2.keras")
      test_loss, test_acc = model.evaluate(test_dataset)
      print(f"Test accuracy: {test_acc:.3f}")

```

63/63 [=====] - 107s 2s/step - loss: 0.6919 - accuracy: 0.9325
Test accuracy: 0.933

```

[48]: conv_base = keras.applications.vgg16.VGG16(
      weights="imagenet",
      include_top=False,
      input_shape=(180, 180, 3))

```

```

[49]: conv_base = keras.applications.vgg16.VGG16(
      weights="imagenet",
      include_top=False)

conv_base.trainable = True
for layer in conv_base.layers[:-4]:
    layer.trainable = False

```

```

[51]: data_augmentation = keras.Sequential(
      [
          layers.RandomFlip("horizontal"),
          layers.RandomRotation(0.1),
          layers.RandomZoom(0.2),
      ]
    )

```

```

inputs = keras.Input(shape=(180, 180, 3))
x = data_augmentation(inputs)
x = keras.applications.vgg16.preprocess_input(x)
x = conv_base(x)
x = layers.Flatten()(x)
x = layers.Dense(256)(x)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs, outputs)
model.compile(loss="binary_crossentropy",
              optimizer=keras.optimizers.RMSprop(learning_rate=1e-5),
              metrics=["accuracy"])

callbacks = [
    keras.callbacks.ModelCheckpoint(
        filepath="fine_tuning3.keras",
        save_best_only=True,
        monitor="val_loss")
]
history = model.fit(
    train_dataset,
    epochs=5,
    validation_data=validation_dataset,
    callbacks=callbacks)

```

63/63 [=====] - 192s 3s/step - loss: 3.8077 - accuracy: 0.7520 - val_loss: 0.8551 - val_accuracy: 0.9150

```

[53]: model = keras.models.load_model("fine_tuning3.keras")
test_loss, test_acc = model.evaluate(test_dataset)
print(f"Test accuracy: {test_acc:.3f}")

```

63/63 [=====] - 107s 2s/step - loss: 0.8319 - accuracy: 0.9205
Test accuracy: 0.920